

Agentic AI Readiness

What was built, what it changes, what remains, and how to measure whether it worked. Covers all three storefronts.

DATE	APPLIES TO	SPEC COMPLETE	DEPLOYED
12 August 2026	Solana · BBQ · OKO	97.5%	Not yet

1. Summary

The site is now readable, queryable and callable by AI agents. Scored against the twenty deliverables in the client specification:

19 DELIVERED				1
20 deliverables	19.5 weighted score	97.5% complete	1 scoped out · nothing blocked	

Updated 12 August. The return policy was the one item still waiting on the client, and it turned out not to need them: the return window, the restocking fee and who pays return postage were already published on /return-policy. It is now in the structured data, which moves the score from 95% to 97.5% and leaves **nothing blocked**.

The remaining half point is agent-callable cart and holds, excluded from this phase on purpose. Scoring is unchanged from the original method — delivered counts 1, partial or blocked counts a half — so the figures stay comparable rather than being flattered by a change of rules.

The headline result. Before this work, every product listing page returned **zero products** to anything that does not run JavaScript, and the storefront rendered a **blank page** entirely — markup was present in the HTML but React streamed it into a hidden element only client-side JavaScript could reveal. Both are fixed on all three brands.

MEASURE	BEFORE	NOW
Products readable on a category page, no JavaScript	0	30
Products readable on brand / main-nav / search pages	0	30
Page renders at all with JavaScript disabled	blank	full
Structured-data types emitted	1	6
AI crawlers with an explicit policy	0	21
Machine-callable catalogue tools	0	2

2. What was built

Readable

- Listing pages server-render their products
- Whole storefront renders without JavaScript
- `Organization` , `WebSite` , `Product` , `BreadcrumbList` , `ItemList` , `FAQPage`
- Product schema carries GTIN, part number, condition, price validity and itemised specifications
- Return policy published as data — 30 days, 20% restocking fee, return postage on the customer
- Semantic landmarks and aria-labels on product pages, which had none

Discoverable

- `llms.txt` — site index for agents
- `llms-full.txt` — policies, taxonomy, brands
- 21 AI crawlers named explicitly in `robots.txt`
- Both files regenerate on their own; no deploy needed

Queryable

- `/api/catalog/search` — text, brand, category, price filters
- `/api/catalog/product/{handle}` — full detail with specifications
- `/openapi.json` — machine-readable API description
- Rate limited: 429 with `Retry-After`, so agents back off gracefully

Callable

- MCP server — assistants query the catalogue as a native tool
- Two tools: search products, get product
- Anti-hallucination rules carried in the tool descriptions
- Tool failures reported honestly rather than invented around

Bugs found and fixed along the way

- A product prefetch on brand and collection pages that **threw on every single request** and discarded a wasted Elasticsearch call. Running undetected in production.
- The generated `llms.txt` shipping with its category list missing entirely — correct in development, empty in the real build.
- Three catalogue reads that silently dropped every consuming route out of static rendering, filling build logs with errors and costing page speed.
- A separate homepage crash on BBQ, unrelated to this work, diagnosed and fixed during the same period.

3. Client decisions applied

QUESTION	ANSWER	WHAT WE DID
Shipping	Calculated per order at checkout	Shipping details deliberately left unpublished. Schema expects a stated rate or a free-shipping threshold; we have neither, and publishing one would be a promise checkout cannot keep.
Inventory	Orders are taken and sourced	No stock feed to build. Agents are now told plainly that availability means <i>purchasable</i> , not on a shelf, and must not claim physical stock or quote dispatch times.
AI training crawlers	Allow search, block training	9 search crawlers allowed, 12 training crawlers blocked. Visibility in AI answers is unaffected — vendors run these as separate bots. Reversible in one line.

Raised for a different owner. The storefront currently advertises *both* "free shipping on orders \$499+" and "free shipping on orders over \$1,999", in different components. Those contradict each other and the calculated-at-checkout model. This is storefront copy rather than structured data, but it is a promise the checkout may not keep.

4. What remains

ITEM	STATUS	NEEDS
Shipping rate in structured data	Needs a fact	Whether checkout really charges free over \$79.99 and \$9.99 below. The storefront copy now says so consistently, but a claim on a page is not a rate checkout will honour, and structured data is read as a commitment. One confirmation switches it on.
Agent-callable cart and holds	Scoped out	Deliberately excluded. Letting agents create carts and hold stock brings authentication, abuse prevention and revenue risk. Price separately.
Clear duplicated social accounts	Hosting	BBQ and OKO still publish Solana's Facebook and Pinterest, telling agents all three sites are one business. Config change, not code.
Confirm bot protection	Hosting	Hosting-level bot protection can silently block AI crawlers and undo the crawler policy. Needs checking on all three projects.
Deploy	Pending	None of this is live. Three separate deployments, all three must be released.

5. How to measure whether it worked

Immediately — the same day it deploys

```
curl -s $SITE/robots.txt | grep -c '^User-Agent:'      # expect 22
curl -s $SITE/lms.txt | head -1                      # store name
curl -s "$SITE/api/catalog/search?q=grill&limit=1"   # JSON, not HTML
curl -s $SITE/category/grills-and-smokers \
  | grep -c 'href=".*product/'                       # expect 30, was 0
```

Then paste any product URL into **validator.schema.org**. It should report `Product`, `BreadcrumbList` and `FAQPage` with no errors.

Weeks 1–4 — server logs

Filter access logs by user-agent for `OAI-SearchBot`, `PerplexityBot` and `Claude-SearchBot`. Rising hits mean AI crawlers are now fetching the catalogue. **This is the most direct signal available**, and it did not exist before.

Weeks 2–8 — the outcome that matters

Ask ChatGPT, Claude and Perplexity: "Where can I buy a 4-burner built-in gas grill under \$3,000?" Do the brands appear? Repeat monthly and keep the answers. Crude, but it is the result the business actually cares about.

Ongoing — Search Console

The server-rendering fix is ordinary technical SEO. Expect impressions on category and brand pages to move independently of anything AI-related, because those pages were previously empty to any crawler that does not execute JavaScript.

Set this expectation before anyone asks for a dashboard. AI referral traffic is largely unattributable — most assistants send visitors through a plain link with no distinguishing referrer. Judge this work on *crawler activity* and *appearing in answers*, not on a single number in analytics. Anyone promising a clean attribution figure for AI traffic is overselling.

6. Basis

- Verified on local production builds of all three brands, 7–11 August 2026. Every brand builds clean at 340/340 pages.
- Every figure was measured, not estimated — including loading pages in a real browser with JavaScript disabled, which is what exposed the blank-page fault that source-level checks had missed.
- Development and production rendering can differ. Confirm the checks in section 5 against the live sites before treating any of this as a release record.
- Engineering detail is maintained alongside the code at `docs/agent-ai-readiness.md`.