

Week of 5–11 August 2026

Agentic AI readiness taken from audit to shipped, a multi-store admin layer, and one production outage resolved. Covers all three storefronts.

PERIOD	APPLIES TO	COMMITTS	SPEC COMPLETE
Wed 5 – Tue 11 Aug	Solana · BBQ · OKO	23	48% → 95%

1. Summary

Three things happened this week. The agentic-AI programme went from a baseline audit on Friday to a delivered feature set by Tuesday. A store-scoped admin layer landed on Thursday, giving each brand its own configuration without a code deploy. And a live outage on **bbqgrilloutlet.com** — every page showing an application error — was diagnosed and fixed.

23 COMMITTS	112 FILES CHANGED	+8,282 LINES ADDED	62 NEW FILES
----------------	----------------------	-----------------------	-----------------

Scored against the twenty deliverables in the client specification, completion moved from **48%** (first scored Friday) to **95%**:

18 DELIVERED			1	1
--------------	--	--	---	---

20 deliverables 95% complete 1 partial · 1 blocked on the client

The headline result. Before this week, every product listing page returned **zero products** to anything that does not run JavaScript, and the storefront rendered a **blank page** entirely. The markup was in the HTML, but React streamed it into a hidden element that only client-side JavaScript revealed — so every earlier measurement of "the content is in the source" was technically true and practically worthless. Fixed on all three brands.

MEASURE	BEFORE	NOW
Products readable on a category page, no JavaScript	0	30
Products readable on brand / main-nav / search pages	0	30
Page renders at all with JavaScript disabled	blank	full
Structured-data types emitted	1	6
AI crawlers with an explicit policy	0	21
Machine-callable catalogue tools	0	2
Rate-limited public endpoints	0	5
Build warnings bailing routes out of static rendering	many	0

2. What happened, day by day

- Thursday 6 August** ADMIN & MULTI-STORE
Store-scoped configuration: `STORE_ID` became the single bootstrap identity, moving per-brand config out of environment variables and into Redis. Three new admin screens — Store Settings, Page SEO (wired into 19 routes) and XML Feeds with on-demand regeneration. The 1,667-line menu editor was split into per-panel routes, fixing five bugs found along the way, including FAQs being silently lost on save. Cache keys were scoped by store after BBQ was caught reporting OKO's merchant feed. Per-brand dev/build scripts added.
- Friday 7 August** AUDIT & FOUNDATIONS
Baseline audit measured every route: only the product page was machine-readable. Shared schema.org builders landed so the three themes cannot drift, along with `Organization`, `WebSite`, `ItemList`, `BreadcrumbList` and `FAQPage` markup, an explicit AI-crawler policy and `/llms.txt`. A single clear-cache control replaced a "clear all" that had been busting 6 of 13 caches. A spike proved listing grids could be server-rendered. Readiness report issued at 48%.

Monday 10 August MAKING IT REAL

The blank-page-without-JavaScript failure was found and fixed — the single most consequential change of the week. Product pages gained semantic landmarks and labels; the category index gained structured data; `/llms-full.txt` was added. Product specifications became machine-readable (24 typed attributes per product). Rate limiting shipped ahead of the public API, not after it. Report refreshed to 60%.

Tuesday 11 August QUERY LAYER & OUTAGE

Started with the BBQ outage fix. Then a performance fix removing every static-rendering bailout, the public catalogue API with its OpenAPI spec, and an MCP server letting assistants call the catalogue directly. Closed with three client decisions applied and the final report issued at 95%.

3. Agentic AI — what an agent can now do

Read the site - Listing pages server-render their products - Whole storefront renders without JavaScript - Six schema.org types across the site - Product schema carries GTIN, part number, condition, price validity and itemised specifications - Five labelled landmarks per product page	Query the catalogue <code>GET /api/catalog/search</code> — free text plus brand, category and price filters <code>GET /api/catalog/product/{handle}</code> — full detail including specifications <code>GET /openapi.json</code> — machine-readable contract <code>POST /api/mcp</code> — JSON-RPC, two tools, so an assistant calls rather than browses
---	---

The practical difference: a question like "4-burner built-in gas grill under \$3,000" was previously unanswerable without rendering and scraping pages. It is now a single tool call — verified returning 2,467 matches with real titles and prices, and a handle from those results resolving to 19 typed specifications.

Protection shipped alongside

There was previously **no throttling of any kind on any endpoint**, while two of them proxied an unauthenticated Elasticsearch query. That had to land before a public catalogue API, not after it.

ENDPOINT	LIMIT	BEHAVIOUR
<code>/api/es/searchkit</code>	120 / min	Every response carries <code>RateLimit-*</code> headers so a client can pace itself; exhaustion answers <code>429</code> with <code>Retry-After</code> . Redis-backed, fails open, internal page rendering exempt.
<code>/api/es/shopify/search</code>	120 / min	
<code>/api/popular_searches</code>	300 / min	

Why a 429 and not a CAPTCHA. A well-behaved agent reads `Retry-After` and backs off. A CAPTCHA makes it give up and recommend a competitor instead.

4. Client decisions applied (11 August)

QUESTION	ANSWER	WHAT WE DID
Shipping rates	Calculated per order at checkout	Shipping is deliberately <i>not</i> asserted in structured data — schema.org expects a stated rate or a free-shipping threshold, and publishing one would be a promise checkout cannot keep.
Inventory feed	None exists, none planned; orders are sourced	No code change needed. Agents are now told plainly that availability means <i>purchasable</i> , not on a shelf, and must not quote a dispatch time.
AI crawler policy	Allow search, block training	9 search crawlers allowed, 12 training crawlers blocked. Visibility in AI shopping answers is unaffected — vendors run these as separate bots precisely so this split works. Reversible by moving one list.

5. Problems found and fixed

Most of these were pre-existing and silent. They are listed because each one was costing something before it was found.

ISSUE	IMPACT	STATUS
BBQ homepage crash. A blog export was being removed by the build optimiser while a reference to it remained	bbqgrilloutlet.com served an application error on <i>every page</i> . OKO carried the same latent fault	Fixed
Storefront invisible without JavaScript. The entire site sat inside a boundary React only revealed on hydration	Blank page to any crawler or agent that does not execute JavaScript. True before this week's work as well	Fixed
Product prefetch had never worked. A response was being read one level off, so it threw on every request and the error was swallowed	Zero server-rendered products on brand pages, plus a wasted Elasticsearch round-trip on <i>every listing render</i>	Fixed

Routes silently bailing out of static rendering. Three catalogue reads conflicted with their pages' caching	Pages shipped statically with empty data while the build log filled with warnings. This is how <code>11ms.txt</code> shipped with no categories	Fixed
"Clear all caches" cleared 6 of 13. The list had drifted from the code	Product data, FAQs, reviews and store settings stayed stale for up to 24 hours after an operator cleared the cache	Fixed
Cache entries leaking across brands. Keys were not scoped by store	BBQ reported OKO's merchant feed during testing	Fixed

6. What remains

Nothing here is confirmed live yet. All work is pushed to both remotes, but each of the three Vercel projects deploys independently and none has been verified in production. That is the immediate next action, and the checks below are the ones a build log will not catch.

ITEM	WHY IT MATTERS	OWNER
Deploy and verify all three brands	Confirm each Vercel project builds from the branch that actually carries this work, then load a category page <i>with JavaScript disabled</i> — the one check no script covers and the exact failure this release fixes	Us
Check Vercel bot protection	There is no Cloudflare or WAF in front of these sites, so bot protection is the only thing that could silently undo the crawler policy	Us
BBQ and OKO social accounts	Both still point at Solana's accounts in configuration. Now that these feed structured data, they actively tell agents the three brands are one entity	Us
Confirm the return policy	The storefront says "30-Day Returns" but the window, who pays return postage and any restocking fee are unconfirmed. Publishing a policy that contradicts the real one is what gets a merchant penalised, so this is left empty until confirmed — the one item still blocked	Client
Free-shipping copy contradiction	The storefront advertises both "\$499+" and "over \$1,999" in different places, and both contradict shipping being calculated at checkout. A copy decision, not a code change	Client
Agentic checkout protocols	Deliberately not built. The payment stack is Braintree and the leading protocol is Stripe-first, so this is a payment-provider question before it is an engineering one. Worth watching, not worth speculating on	Later

7. How it was verified

Every change this week was checked against a running build of all three brands rather than reasoned about. Representative evidence:

- All three brands build **340/340 pages, exit 0**, with zero static-rendering bailouts — down from many at the start of the week
- Category page with JavaScript disabled: **60 product links and 66 prices** on each brand, against a blank page before
- Rate limiter, 340 concurrent requests against a 300/min bucket: exactly **300 accepted, 40 refused** with a correct retry hint, while internal page rendering was unaffected
- Catalogue API and MCP server exercised end to end against a live server, including the failure paths — unknown product, unknown method, backend unreachable
- Structured data confirmed present and intact on product, category and brand pages

Reference. Engineering detail lives in `docs/agentic-ai-readiness.md` ; the full stakeholder report is `AGENTIC-AI-READINESS-REPORT.pdf` at the project root; per-brand release steps are in `docs/deploy-agentic-ai-2026-08-07.md` .